Title:	Methods and Tools for Software Engineering
Course ID:	ECE 650
Term:	Winter 2021
WWW:	`https://ece.uwaterloo.ca/~rbabaeec/ece650/w21`
Campuswire	`https://campuswire.com/p/G9003B726` Enter 3253 when prompted
Lectures:	Wednesday, 14:30 – 17:20 EST
Instructor:	Reza Babaee, `rbabaeec@uwaterloo.ca`
TAs:	An Qi Zhang, `an.zhang@uwaterloo.ca`
	Arshdeep Kaur Bal, `arshdeep.kaur.bal@uwaterloo.ca`

Final Course Project – Due April 21, 2021, 23:59 EST

Cloning

The repository for this assignment is available at `gitlab`. Run the following to clone your repository.

```
git clone ist-git@git.uwaterloo.ca:ece650-1211/project/GROUPNAME.git
```

where `GROUPNAME` is the name you chose for your group. Please note since `gitlab` does not accept - (dash) in the repository names, it is replaced with _ (underline). The project *can* be done in groups of 2.

Introduction

For the project you will need to:
- Augment your code from Assignment 4 in the way that is described below.
- Quantitatively analyze your software for various kinds of inputs.
- Write a brief report (≈ 5 pages, 11 pt font, reasonable margins) with your analysis. Your report should be typeset in LaTeX, and must be in PDF.

You should augment your code from Assignment 4 in the following ways.
- Make it multithreaded. You should have at least 4 threads: one for I/O, and one each for the different approaches to solve the minimum vertex cover problem. You may find the examples on Learn, under concurrency-demo, useful.
- Implement the following two additional ways to solve MIN-VERTEX-COVER, in addition to the REDUCTION-TO-CNF-SAT approach you had in Assignment 4. (We will call your approach from Assignment 4, CNF-SAT-VC.)

  1. Pick a vertex of highest degree (most incident edges). Add it to your vertex cover and throw away all edges incident on that vertex. Repeat till no edges remain. We will call this algorithm APPROX-VC-1.

  2. Pick an edge $\langle u, v \rangle$, and add both u and v to your vertex cover. Throw away all edges attached to u and v. Repeat till no edges remain. We will call this algorithm APPROX-VC-2.

Inputs

As input, use the output of `/home/agurfink/ece650/graphGen/graphGen` on `eceubuntu`. That program generates graphs with the same number of edges for a particular number of vertices, but not necessarily the same edges. Note that you can store its output in a file and use the file on other machines.

Output

Given a graph as input, your program should output the vertex cover computed by each approach in sorted order. That is, give the following input:

```
V 5
E {<2,1>,<2,0>,<2,3>,<1,4>,<4,3>}
```

The output from your program should be:

```
CNF-SAT-VC: 2,4
APPROX-VC-1: 2,4
APPROX-VC-2: 0,2,3,4
```

That is, the name of the algorithm, followed by a colon ':', a single space, and then the computed result as a sorted sequence of vertices, separated by commas.

Analysis

You should analyze how efficient each approach is, for various inputs. An input is characterized by the number of vertices. "Efficient" is characterized in one of two ways: (1) running time, and (2) approximation ratio. We characterize the approximation ratio as the ratio of the size of the computed vertex cover to the size of an optimal (minimum-sized) vertex cover.

For measuring the running time, use `pthread_getcpuclockid()`. For an example of how it is used, see http://www.kernel.org/doc/man-pages/online/pages/man3/pthread_getcpuclockid.3.html.

For measuring the approximation ratio, compare it to the output of CNF-SAT-VC, which is guaranteed to be optimal.

Your objective is to measure, for various values of $|V|$ (number of vertices), for the graphs generated by `graphGen`, the running time and approximation ratio. You should do this by generating graphs for $|V| \in [5, 50]$ using that program, in increments of 5. That is, graphs with $5, 10, 15, \ldots, 50$ vertices.

You should generate at least 10 graphs for each value for $|V|$, compute the time and approximation ratio for each such graph. You should measure the running time for at least 10 runs of each such graph. Then, you should compute the mean (average) and standard deviation across those 100 runs for each value of $|V|$. For the approximation ratio, if there is any random component (e.g., which edges you choose, for APPROX-VC-2), then you should measure that multiple times as well for each graph.

You might find the optimal approach (CNF-SAT-VC) difficult to scale to large number of vertices. For these instances, you can use a timeout to avoid waiting for its result and produce `CNF-SAT-VC: timeout` in the output.

CNF-SAT-VC can be scaled significantly by improving the encoding. Should you decide to improve the encoding, bonus points will be given for scaling to larger instances. Make sure to describe the improvements you made to the encoding in the report. Points will be deducted for encodings that are scalable but are either incorrect or unexplained.

Report

The main part of your report are graphs (plots) corresponding to the data you generate as described in the "Analysis" section above. One way to show the output is to have two plots: one for running times and the other for approximation ratio. The horizontal axis is the number of vertices.

Figure 1: Example plot, generated using `gnuplot`. The error bars for Approx-2 are not visible because the standard deviation is small.

You should plot the mean for each value of $|V|$ for which you made measurements, and the standard deviation as an errorbar [1]. An example of a possible plot is shown in Figure 1.

The remainder of your report should be reasoning about your plots. That is, you should explain why your plots look the way they do. For example, if there is a "spike" in the approximation ratio for some value of $|V|$ for one of the approaches, you should explain why there is such a spike. You should also explain apparent trends. For example, if, for one of the approaches, the running time seems to increase linearly with $|V|$, you should reason about why that is happening.

Marking

We will mark by: (1) Trying some inputs and checking your output, (2) inspecting your code to make sure that you are using pthreads correctly, and, (3) reading your report.
- Marking script for compile/make etc. fails: automatic 0
- Your program runs, awaits input and does not crash on input: + 20
- Correctly implemented 2 new algorithms: + 20 each, total + 40
- Generated plots: + 20
- Report: + 20

CMake

As discussed below under "Submission Instructions", you should use a `CMakeLists.txt` file to build your project. We will build your project using the following sequence:

---

[1]You can use yerrorbar in gnuplot to draw an error bar: `http://gnuplot.sourceforge.net/docs_4.2/node262.html`

3

```
cd PROJECT && mkdir build && cd build && cmake ../ && make
```

where `PROJECT` is the top level directory of your submission. You can *assume* that MiniSat will be placed in directory `PROJECT/minisat` (there is no need to push that into your repository on `gitlab`). Your submission should only include your own code. If your code is not compiled from scratch (i.e., from the C++ sources), you get an automatic 0. Unlike for the assignments, you must create the `CMakeLists.txt` file on your own. You can use a `CMakeLists.txt` file from previous projects or from the course examples provided on GitHub.

Submission Instructions

You should place all your files at the top of a GitHub repository. The repository should contain:

- All your C++ source-code files.
- A `CMakeLists.txt`, that builds your C++ executable `prjece650`.
- A file `user.yml` that includes your name, WatIAM, and student number of all the team members, as well as the hours you put into the project (please add the *hours:* field yourself with the number in front). Note that *WatIAM* is the user name for your Quest account, e.g. `agurfink`, and a *student number* is an 8-digit number, e.g. `20397238`. If you have done the assignment as a group, the information for both members of the group should be included in the `user.yml`.
- A file named "`report.pdf`" with your report.

See `README.md` for any additional information.

The submitted files should be at the top of the `master` branch of your repository.